VoiceOps 2026

Submission verification report | 29 October 2026

SPEECH CHECKS: PASS | PUBLIC-PACKAGE GATE: PASS WITH ONE PRIVATE-INPUT TEST DESELECTED

This report documents measured offline validation and the latest bounded synthetic-audio smoke test. It does not claim a guaranteed competition score, a fix for the historical STT failure, browser microphone validation, or complete production end-to-end verification.

Verification gates

Gate	Observed result
Public-package tests	69 passed, 1 deselected. The deselected test requires eval/hidden_cases.csv, intentionally excluded. The unfiltered suite is not reported as passing.
Python compilation	Passed (exit 0).
Dataset validator	Passed (exit 0).
Gnani synthetic-clip smoke test	STT passed with one request; TTS passed with one request.

Recorded evaluation metrics

Evaluation	Content / decision / checks	Recall
Offline normal (52)	100% / 100%	0.968
Offline ASR-text (52)	100% / 100%	0.968
Offline conversations (23; 96 turns)	23/23 fully correct; 100% checks	0.944
Latest bounded RAG normal	98.1% content; 96.2% decision	0.968
Latest bounded RAG ASR-text	100% content; 94.2% decision	0.968
Latest bounded RAG conversations (96 turns)	100% checks; 23/23 fully correct	0.944

Known retrieval regression

Compared with the verified saved baseline, normal and ASR offline recall declined from 0.989 to 0.968; conversation recall declined from 1.000 to 0.944. Completed rollback experiments did not recover the baseline. Accuracy and conversation correctness remain as listed above. A previous ASR RAG attempt ended in provider transport failure and offline fallback; it is invalid RAG evidence.

Speech integration and limitations

Bounded test on the bundled prerecorded synthetic demo; provider retries were disabled.

Observed speech results

Operation	Result	Requests
STT	PASS: provider returned a nonempty transcript. Transcript content was not printed or retained in this report.	1
TTS	PASS: generated WAV was readable and nonempty (2.24 seconds; 215,084 bytes).	1

The STT input was normalized to mono 16-bit PCM at 16 kHz; measured duration was 12 seconds and RMS was 3912.6. Installed SDK was gnani-vachana 0.8.2. The historical STT failure was not reproduced, and its root cause is not established. No claim is made that the earlier failure is fixed.

Scope and remaining limitations

• The smoke check used prerecorded synthetic audio. Browser microphone behavior, customer audio, and full production end-to-end operation were not tested.

• One hidden-case-dependent pytest test was deselected because its private input is excluded. Do not describe the unfiltered test suite as passing.

• Embedding retrieval was not validated with a confirmed embedding model. The diagnostic held-out set was highly templated and is weak evidence of general retrieval quality.

• The candidate retains the measured retrieval-recall reduction. No offline evaluations were rerun for this speech update; previously recorded results are reproduced above.

• Credentials, datasets, hidden labels, raw evaluation results, caches, request ledgers, temporary audio, logs, and backups are excluded from the submission package.